

Constraint Satisfaction Problems (CSPs)

This lecture topic (two lectures)
Chapter 6.1 – 6.4, except 6.3.3

Next lecture topic (two lectures)
Chapter 7.1 – 7.5

(Please read lecture topic material before and after
each lecture on that topic)

Outline

- What is a CSP
- Backtracking for CSP
- Local search for CSPs
- ~~(Removed) Problem structure and decomposition~~

You Will Be Expected to Know

- Basic definitions (section 6.1)
- Node consistency, arc consistency, path consistency (6.2)
- Backtracking search (6.3)
- Variable and value ordering: minimum-remaining values, degree heuristic, least-constraining-value (6.3.1)
- Forward checking (6.3.2)
- Local search for CSPs: min-conflict heuristic (6.4)

Constraint Satisfaction Problems

- What is a CSP?
 - Finite set of variables X_1, X_2, \dots, X_n
 - Nonempty domain of possible values for each variable D_1, D_2, \dots, D_n
 - Finite set of constraints C_1, C_2, \dots, C_m
 - Each constraint C_i limits the values that variables can take,
 - e.g., $X_1 \neq X_2$
 - Each constraint C_i is a pair $\langle \text{scope}, \text{relation} \rangle$
 - Scope = Tuple of variables that participate in the constraint.
 - Relation = List of allowed combinations of variable values.
May be an explicit list of allowed combinations.
May be an abstract relation allowing membership testing and listing.
- CSP benefits
 - Standard representation pattern
 - Generic goal and successor functions
 - Generic heuristics (no domain specific expertise).

Sudoku as a Constraint Satisfaction Problem (CSP)

- Variables: 81 variables

- A1, A2, A3, ..., I7, I8, I9
- Letters index rows, top to bottom
- Digits index columns, left to right

	1	2	3	4	5	6	7	8	9
A		6		1		4		5	
B			8	3		5	6		
C	2								1
D	8			4		7			6
E			6				3		
F	7			9		1			4
G	5								2
H			7	2		6	9		
I		4		5		8		7	

- Domains: The nine positive digits

- $A1 \in \{1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Etc.

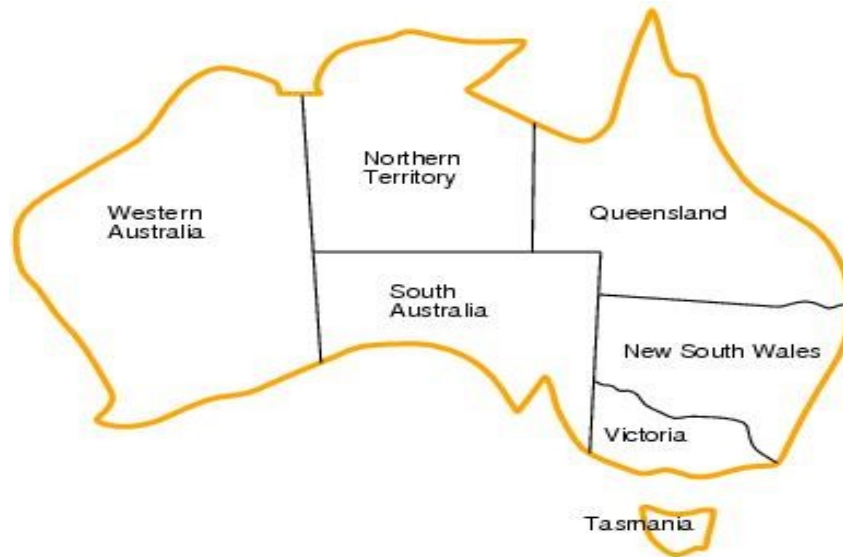
- Constraints: 27 *Alldiff* constraints

- *Alldiff*(A1, A2, A3, A4, A5, A6, A7, A8, A9)
- Etc.

CSPs --- what is a solution?

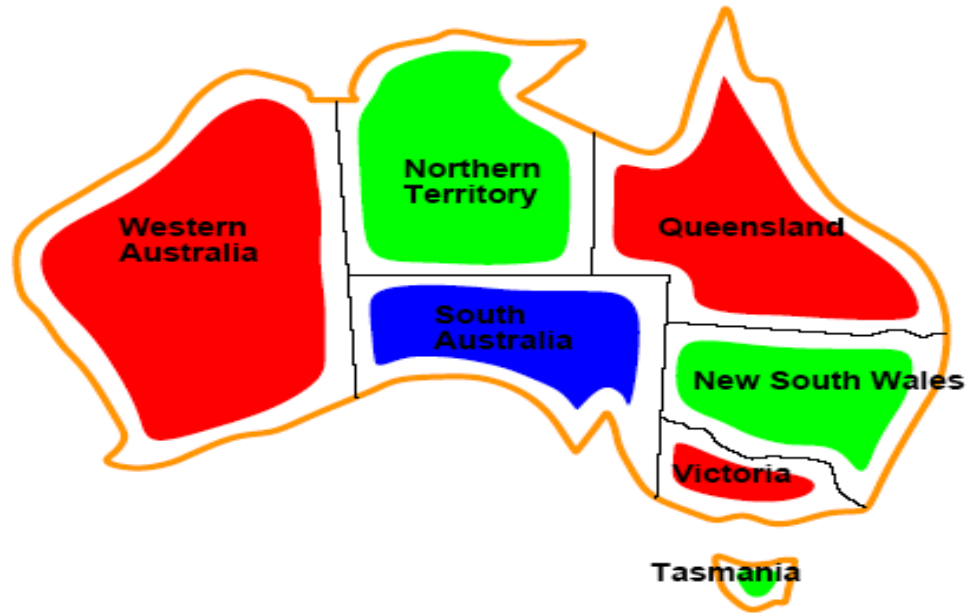
- A *state* is an *assignment* of values to some or all variables.
 - An assignment is *complete* when every variable has a value.
 - An assignment is *partial* when some variables have no values.
- **Consistent assignment**
 - assignment does not violate the constraints
- A **solution** to a CSP is a complete and consistent assignment.
- Some CSPs require a solution that maximizes an *objective function*.
- Examples of Applications:
 - Scheduling the time of observations on the Hubble Space Telescope
 - Airline schedules
 - Cryptography
 - Computer vision -> image interpretation
 - Scheduling your MS or PhD thesis exam 😊

CSP example: map coloring



- Variables: WA, NT, Q, NSW, V, SA, T
- Domains: $D_i = \{red, green, blue\}$
- Constraints: adjacent regions must have different colors.
 - E.g. $WA \neq NT$

CSP example: map coloring



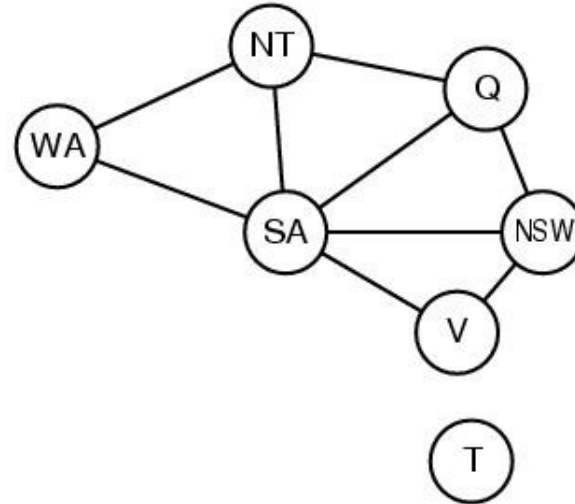
- Solutions are assignments satisfying all constraints, e.g.
 $\{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green\}$

Graph coloring

- More general problem than map coloring
- Planar graph = graph in the 2d-plane with no edge crossings
- Guthrie's conjecture (1852)
 - Every planar graph can be colored with 4 colors or less*
 - Proved (using a computer) in 1977 (Appel and Haken)

Constraint graphs

- Constraint graph:
 - nodes are variables
 - arcs are binary constraints



- Graph can be used to simplify search
 - e.g. Tasmania is an independent subproblem
- (will return to graph structure later)

Varieties of CSPs

- Discrete variables
 - Finite domains; size $d \Rightarrow O(d^n)$ complete assignments.
 - E.g. Boolean CSPs: Boolean satisfiability (NP-complete).
 - Infinite domains (integers, strings, etc.)
 - E.g. job scheduling, variables are start/end days for each job
 - Need a constraint language e.g. $StartJob_1 + 5 \leq StartJob_3$.
 - Infinitely many solutions
 - Linear constraints: solvable
 - Nonlinear: no general algorithm
- Continuous variables
 - e.g. building an airline schedule or class schedule.
 - Linear constraints solvable in polynomial time by LP methods.

Varieties of constraints

- Unary constraints involve a single variable.
 - e.g. $SA \neq green$
- Binary constraints involve pairs of variables.
 - e.g. $SA \neq WA$
- Higher-order constraints involve 3 or more variables.
 - Professors A, B, and C cannot be on a committee together
 - Can always be represented by multiple binary constraints
- Preference (soft constraints)
 - e.g. *red* is better than *green* often can be represented by a cost for each variable assignment
 - combination of optimization with CSPs

CSPs Only Need Binary Constraints!!

- Unary constraints: Just delete values from variable's domain.
- Higher order (3 variables or more): reduce to binary constraints.
- Simple example:
 - Three example variables, X, Y, Z .
 - Domains $D_x = \{1, 2, 3\}$, $D_y = \{1, 2, 3\}$, $D_z = \{1, 2, 3\}$.
 - Constraint $C[X, Y, Z] = \{\mathbf{X+Y=Z}\} = \{(1, 1, 2), (1, 2, 3), (2, 1, 3)\}$.
 - Plus many other variables and constraints elsewhere in the CSP.
- Create a new variable, W , taking values as triples (3-tuples).
- Domain of W is $D_w = \{(1, 1, 2), (1, 2, 3), (2, 1, 3)\}$.
 - D_w is exactly the tuples that satisfy the higher order constraint.
- Create three new constraints:
 - $C[X, W] = \{ [1, (1, 1, 2)], [1, (1, 2, 3)], [2, (2, 1, 3)] \}$.
 - $C[Y, W] = \{ [1, (1, 1, 2)], [2, (1, 2, 3)], [1, (2, 1, 3)] \}$.
 - $C[Z, W] = \{ [2, (1, 1, 2)], [3, (1, 2, 3)], [3, (2, 1, 3)] \}$.
- Other constraints elsewhere involving X, Y , or Z are unaffected.

CSP Example: Cryptarithmic puzzle

$$\begin{array}{r} \text{ T W O} \\ + \text{ T W O} \\ \hline \text{ F O U R} \end{array}$$

Variables: $F T U W R O X_1 X_2 X_3$

Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

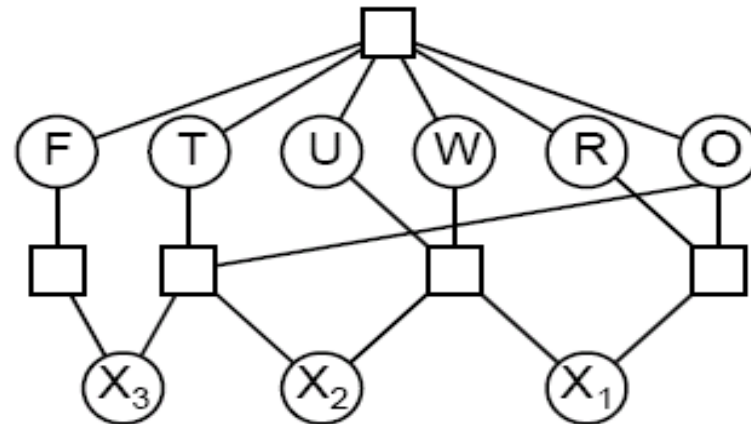
Constraints

$\text{alldiff}(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1$, etc.

CSP Example: Cryptarithmic puzzle

$$\begin{array}{r} \text{ T W O} \\ + \text{ T W O} \\ \hline \text{ F O U R} \end{array}$$



Variables: $F T U W R O X_1 X_2 X_3$

Domains: $\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

Constraints

$alldiff(F, T, U, W, R, O)$

$O + O = R + 10 \cdot X_1$, etc.

CSP as a standard search problem

- A CSP can easily be expressed as a standard search problem.
- Incremental formulation
 - *Initial State*: the empty assignment $\{\}$
 - *Actions* (3^{rd} ed.), *Successor function* (2^{nd} ed.): Assign a value to an unassigned variable provided that it does not violate a constraint
 - *Goal test*: the current assignment is complete
(by construction it is consistent)
 - *Path cost*: constant cost for every step (not really relevant)
- Can also use complete-state formulation
 - Local search techniques (Chapter 4) tend to work well

CSP as a standard search problem

- Solution is found at depth n (if there are n variables).
- Consider using BFS
 - Branching factor b at the top level is nd
 - At next level is $(n-1)d$
 -
- end up with $n!d^n$ leaves even though there are only d^n complete assignments!

Commutativity

- CSPs are commutative.
 - The order of any given set of actions has no effect on the outcome.
 - Example: choose colors for Australian territories one at a time
 - [WA=red then NT=green] same as [NT=green then WA=red]
 - All CSP search algorithms can generate successors by considering assignments for only a single variable at each node in the search tree
 - ⇒ there are d^n leaves
- (will need to figure out later which variable to assign a value to at each node)

Backtracking search

- Similar to Depth-first search, generating children one at a time.
- Chooses values for one variable at a time and backtracks when a variable has no legal values left to assign.
- Uninformed algorithm
 - No good general performance

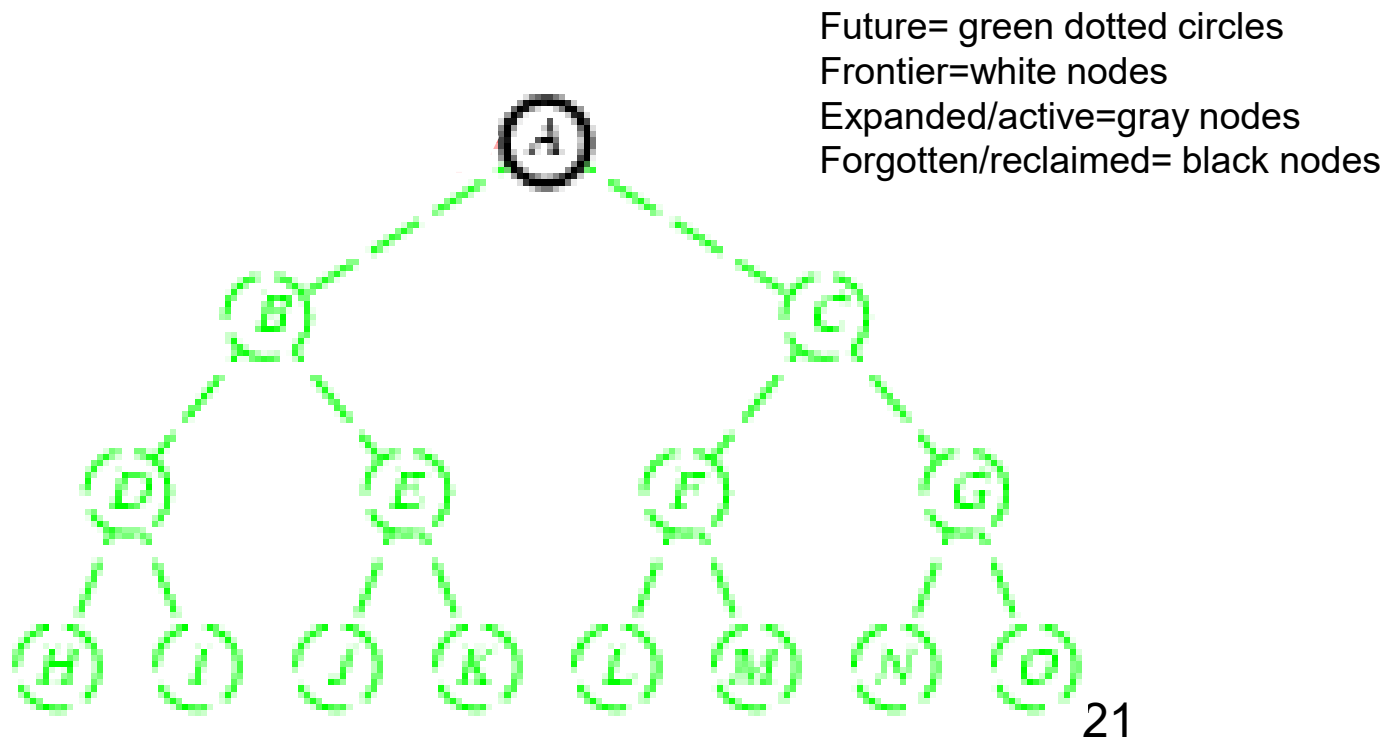
Backtracking search

function BACKTRACKING-SEARCH(*csp*) **return** a solution or failure
 return RECURSIVE-BACKTRACKING($\{\}$, *csp*)

function RECURSIVE-BACKTRACKING(*assignment*, *csp*) **return** a solution or failure
 if *assignment* is complete **then return** *assignment*
 var \leftarrow **SELECT-UNASSIGNED-VARIABLE**(VARIABLES[*csp*],*assignment*,*csp*)
 for each *value* **in** **ORDER-DOMAIN-VALUES**(*var*, *assignment*, *csp*) **do**
 if *value* is consistent with *assignment* according to CONSTRAINTS[*csp*]
 then
 add {*var*=*value*} to *assignment*
 result \leftarrow RECURSIVE-BACKTRACKING(*assignment*, *csp*)
 if *result* \neq failure **then return** *result*
 remove {*var*=*value*} from *assignment*
 return failure

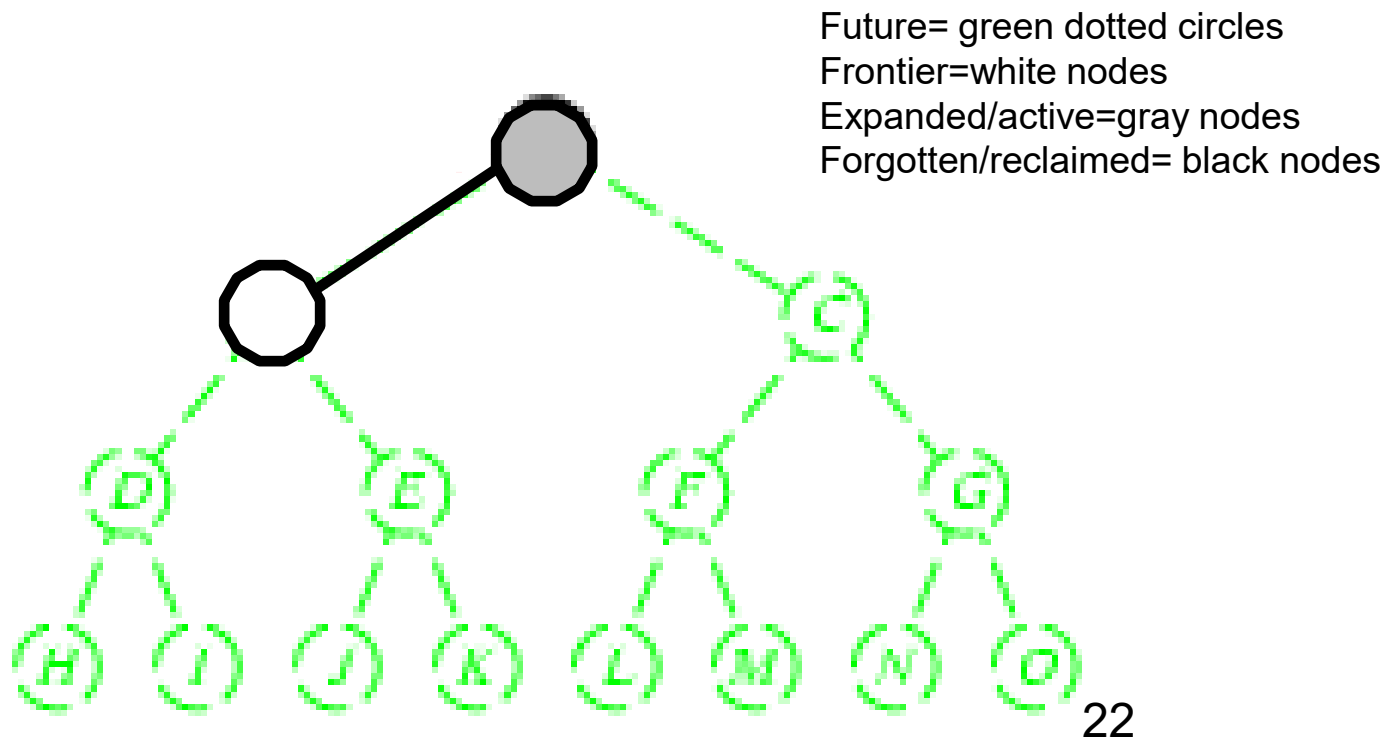
Backtracking search

- Expand *deepest* unexpanded node
- Generate **only one** child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



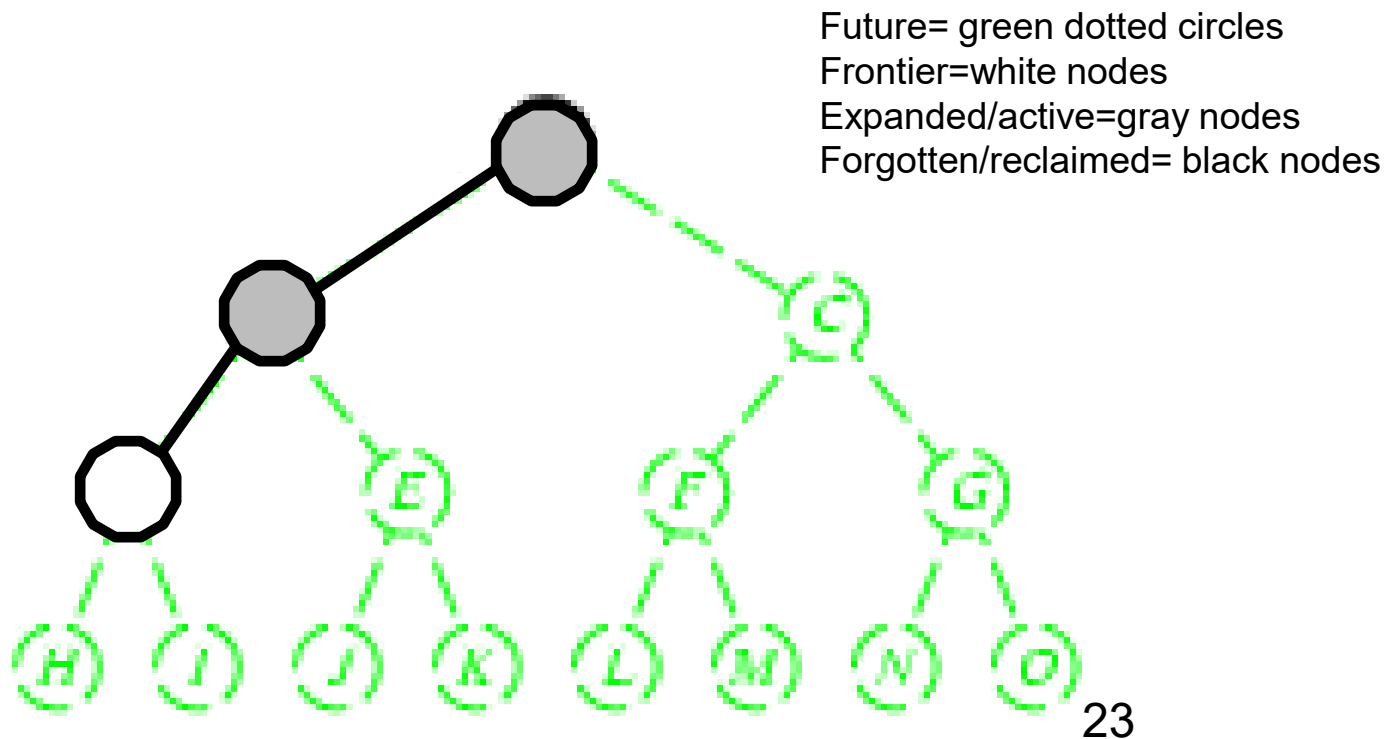
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



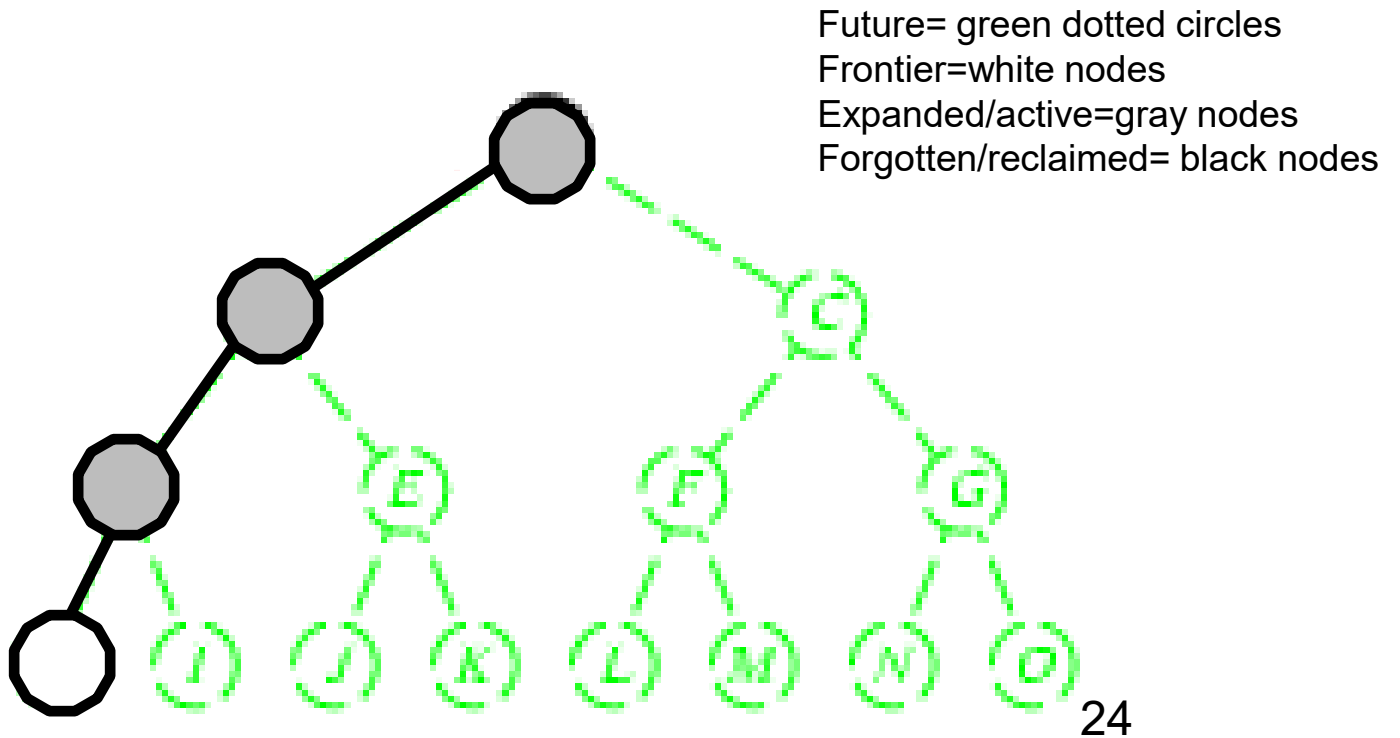
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom




Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



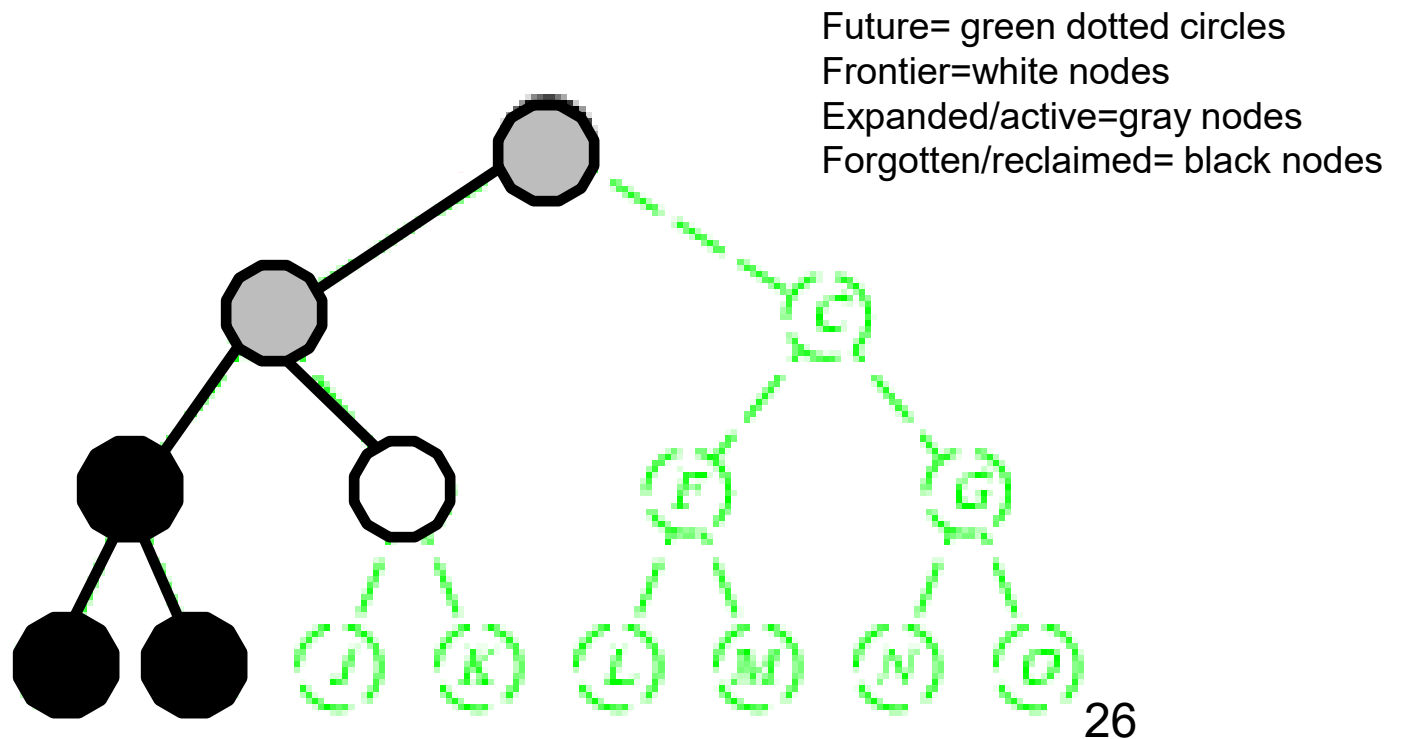
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



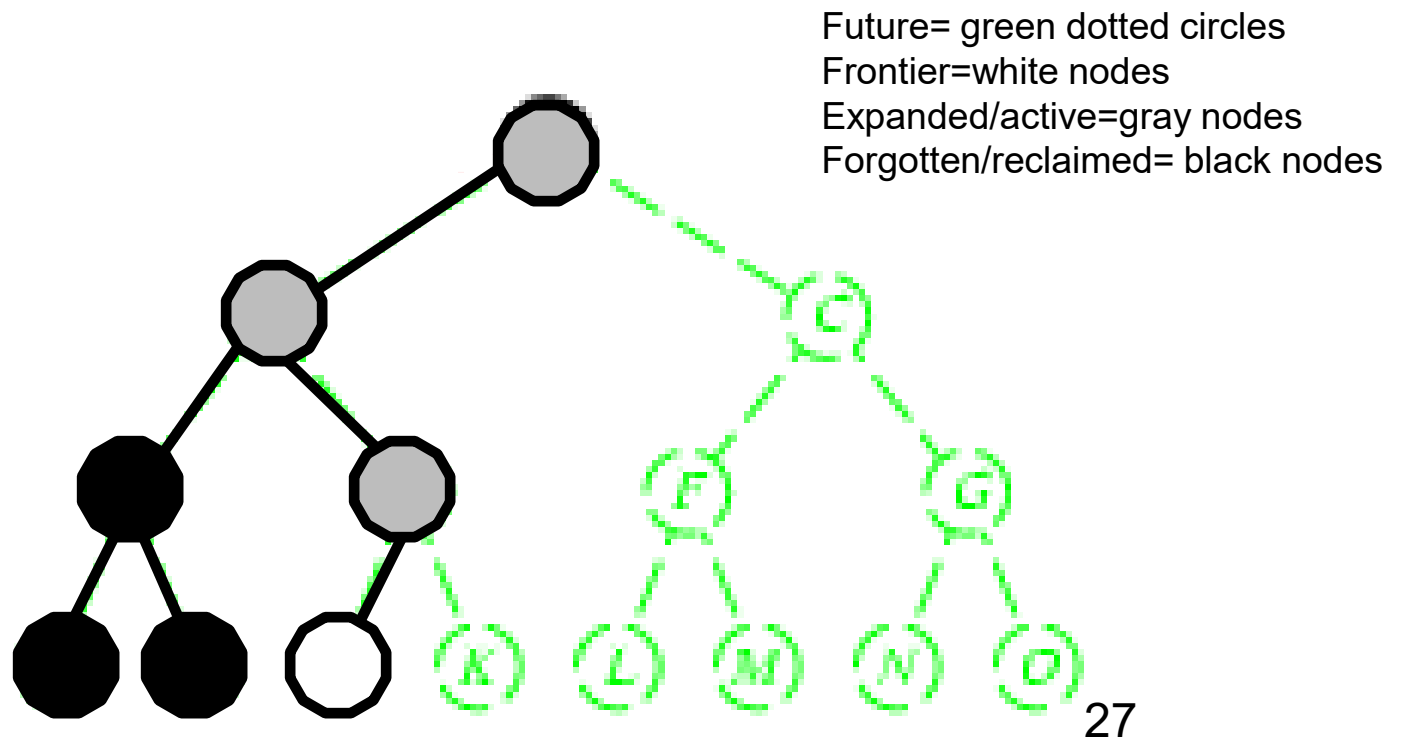
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



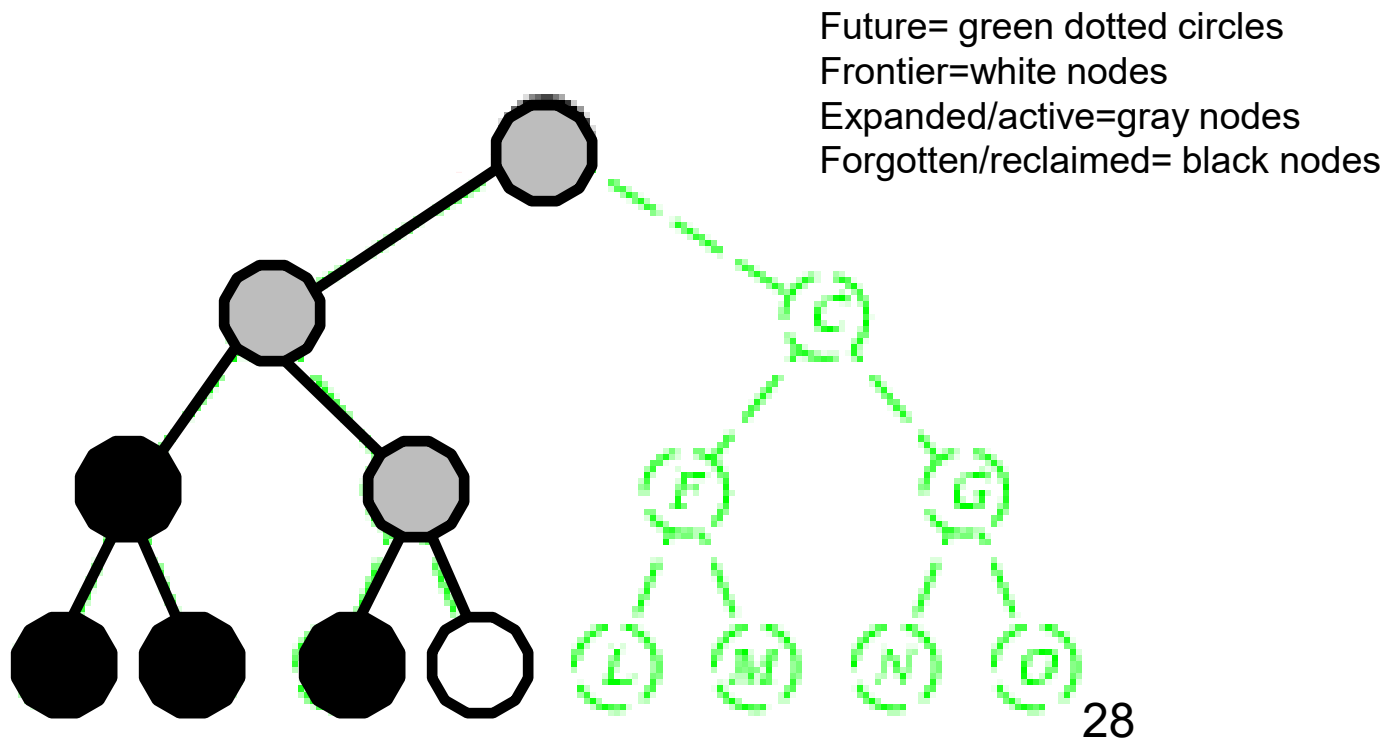
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom




Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



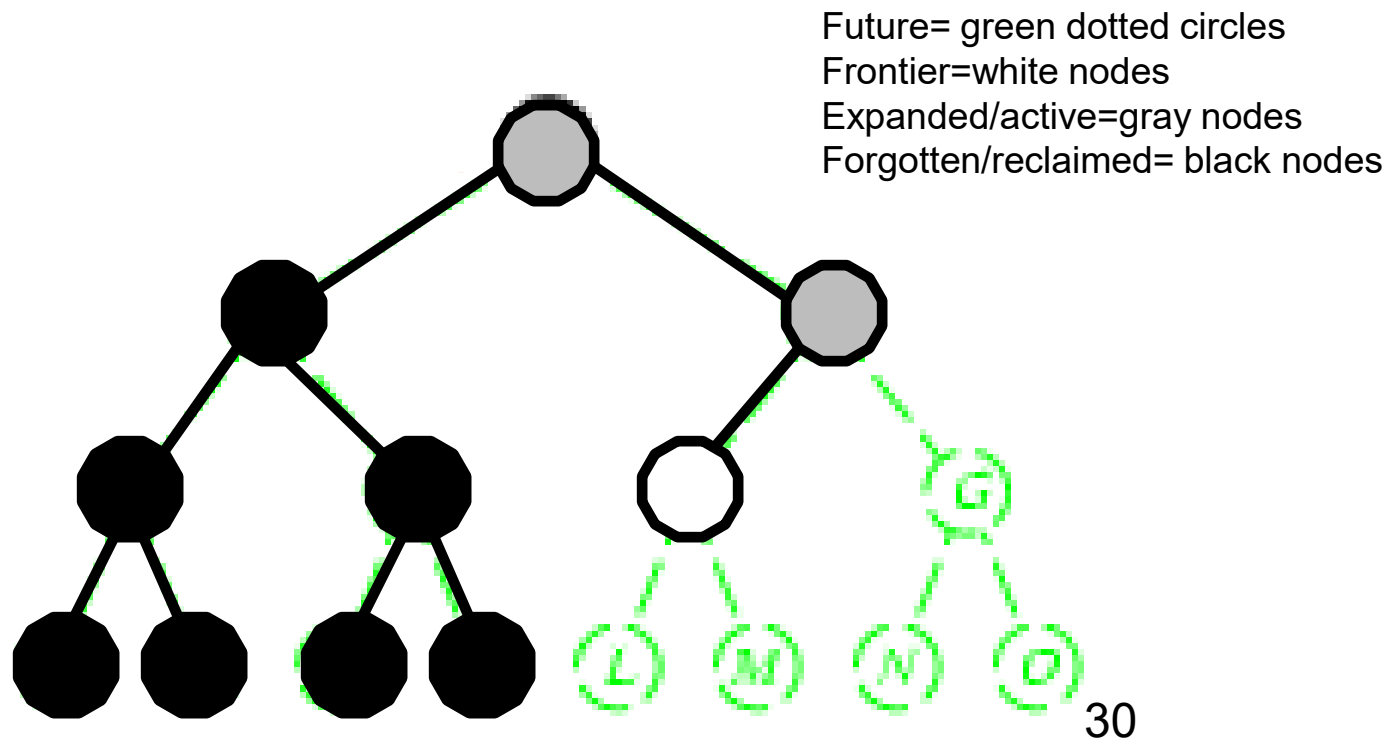
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



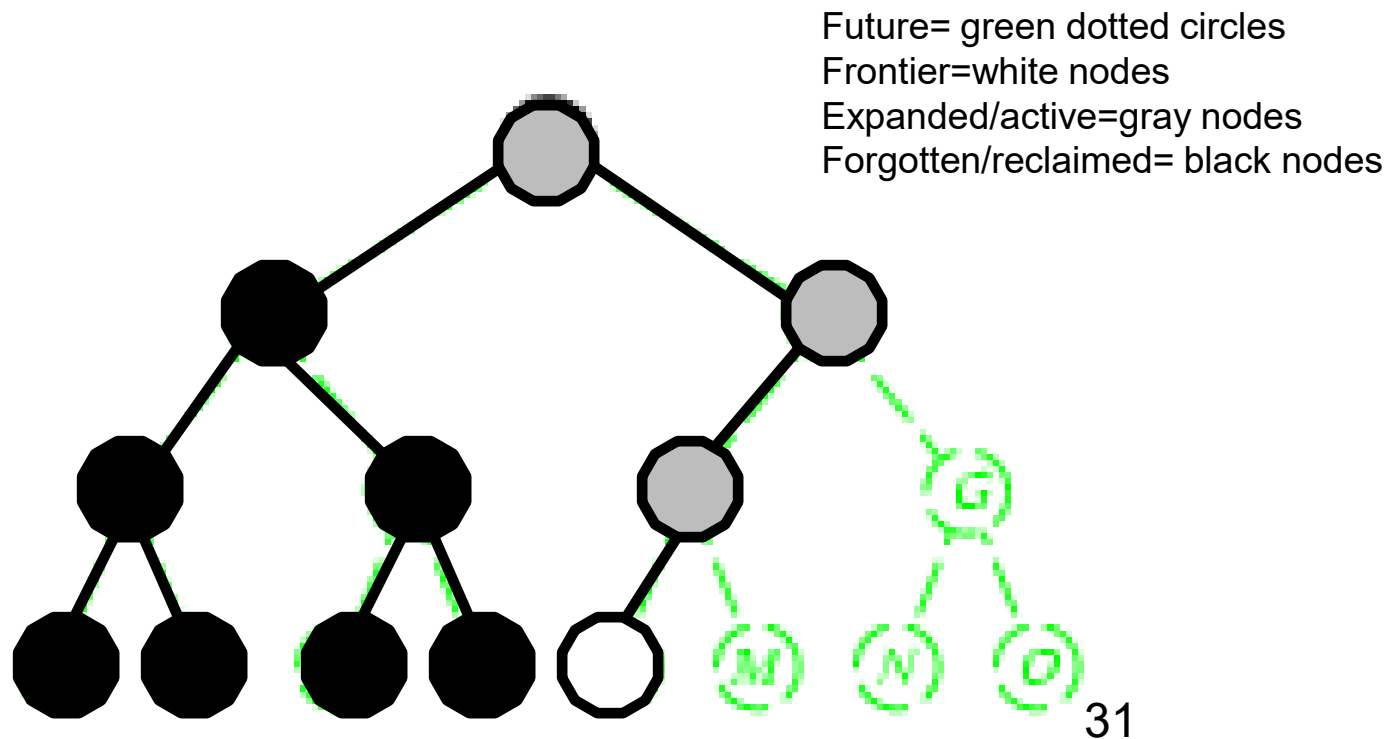
Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom




Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



Backtracking search

- Expand *deepest* unexpanded node
- Generate *only one* child at a time.
- *Goal-Test* when inserted.
 - For CSP, Goal-test at bottom



Backtracking search (Figure 6.5)

function BACKTRACKING-SEARCH(*csp*) **return** a solution or failure
 return RECURSIVE-BACKTRACKING($\{\}$, *csp*)

function RECURSIVE-BACKTRACKING(*assignment*, *csp*) **return** a solution or failure
 if *assignment* is complete **then return** *assignment*
 var \leftarrow **SELECT-UNASSIGNED-VARIABLE(VARIABLES**[*csp*],*assignment*,*csp*)
 for each *value* **in** **ORDER-DOMAIN-VALUES**(*var*, *assignment*, *csp*) **do**
 if *value* is consistent with *assignment* according to CONSTRAINTS[*csp*]
 then
 add {*var*=*value*} to *assignment*
 result \leftarrow RECURSIVE-BACKTRACKING(*assignment*, *csp*)
 if *result* \neq failure **then return** *result*
 remove {*var*=*value*} from *assignment*
 return failure

Comparison of CSP algorithms on different problems

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA	(> 1,000K)	(> 1,000K)	2K	60	64
<i>n</i> -Queens	(> 40,000K)	13,500K	(> 40,000K)	817K	4K
Zebra	3,859K	1K	35K	0.5K	2K
Random 1	415K	3K	26K	2K	
Random 2	942K	27K	77K	15K	

Median number of consistency checks over 5 runs to solve problem

Parentheses -> no solution found

USA: 4 coloring

n-queens: $n = 2$ to 50

Zebra: see exercise 6.7 (3rd ed.); exercise 5.13 (2nd ed.)

Random Binary CSP

(adapted from <http://www.unitime.org/csp.php>)

- A random binary CSP is defined by a four-tuple (n, d, p_1, p_2)
 - n = the number of variables.
 - d = the domain size of each variable.
 - p_1 = probability a constraint exists between two variables.
 - p_2 = probability a pair of values in the domains of two variables connected by a constraint is incompatible.
 - Note that R&N lists compatible pairs of values instead.
 - Equivalent formulations; just take the set complement.
- (n, d, p_1, p_2) are used to generate randomly the binary constraints among the variables.
- The so called model B of Random CSP (n, d, n_1, n_2)
 - $n_1 = p_1 n(n-1)/2$ pairs of variables are randomly and uniformly selected and binary constraints are posted between them.
 - For each constraint, $n_2 = p_2 d^2$ randomly and uniformly selected pairs of values are picked as incompatible.
- The random CSP as an optimization problem (minCSP).
 - Goal is to minimize the total sum of values for all variables.

Improving CSP efficiency

- Previous improvements on uninformed search
 - introduce heuristics
- For CSPs, general-purpose methods can give large gains in speed, e.g.,
 - Which variable should be assigned next?
 - In what order should its values be tried?
 - Can we detect inevitable failure early?
 - Can we take advantage of problem structure?

Note: CSPs are somewhat generic in their formulation, and so the heuristics are more general compared to methods in Chapter 4


Minimum remaining values (MRV)



$var \leftarrow \text{SELECT-UNASSIGNED-VARIABLE}(\text{VARIABLES}[csp], \text{assignment}, csp)$

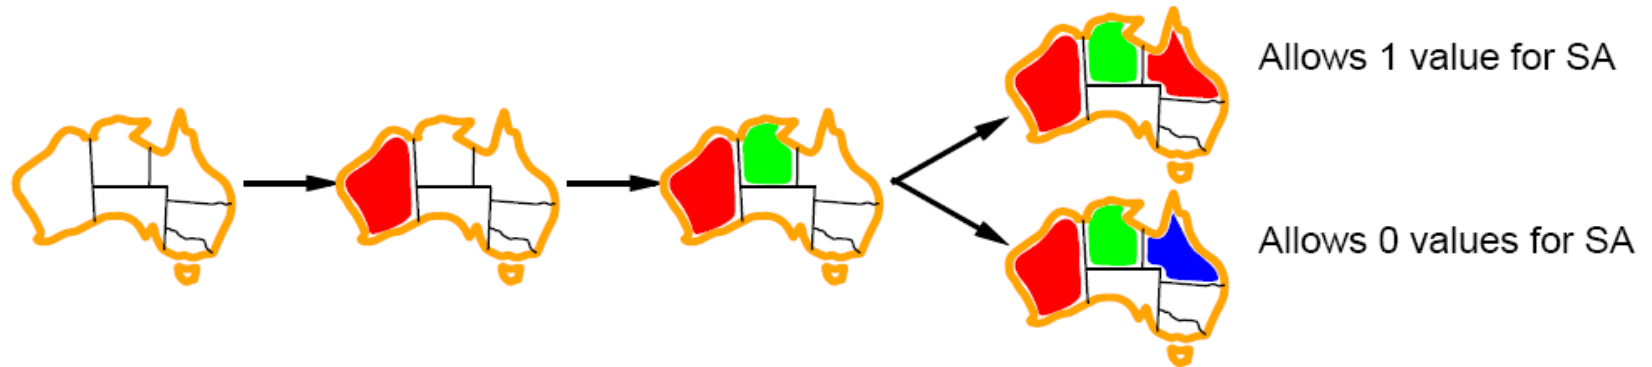
- A.k.a. most constrained variable heuristic
- *Heuristic Rule*: choose variable with the fewest legal moves
 - e.g., will immediately detect failure if X has no legal values

Degree heuristic for the initial variable



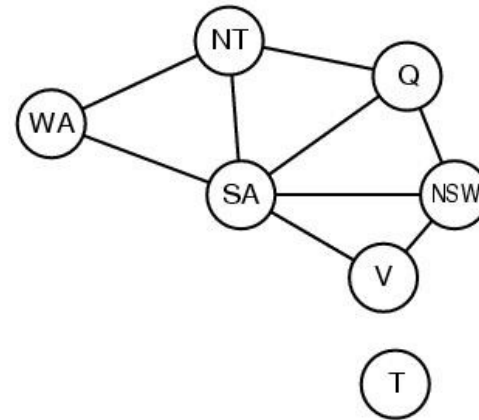
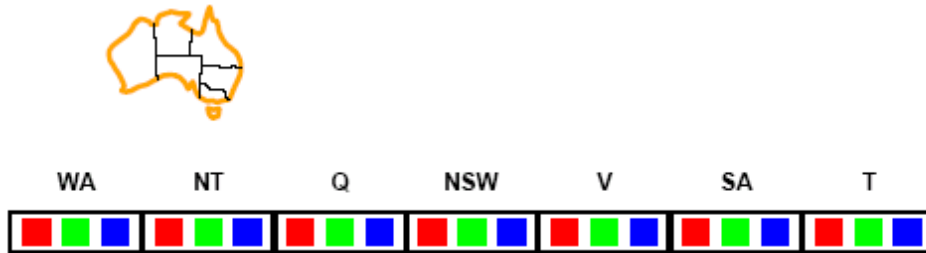
- *Heuristic Rule*: select variable that is involved in the largest number of constraints on other unassigned variables.
- Degree heuristic can be useful as a tie breaker.
- *In what order should a variable's values be tried?*

Least constraining value for value-ordering



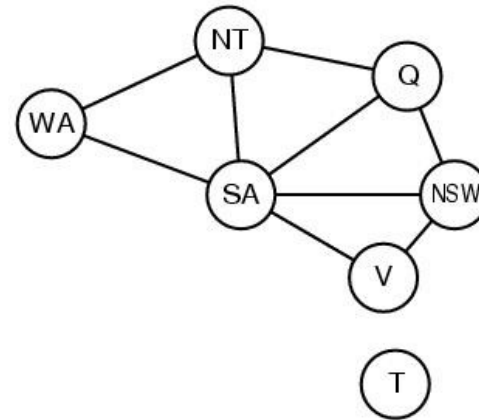
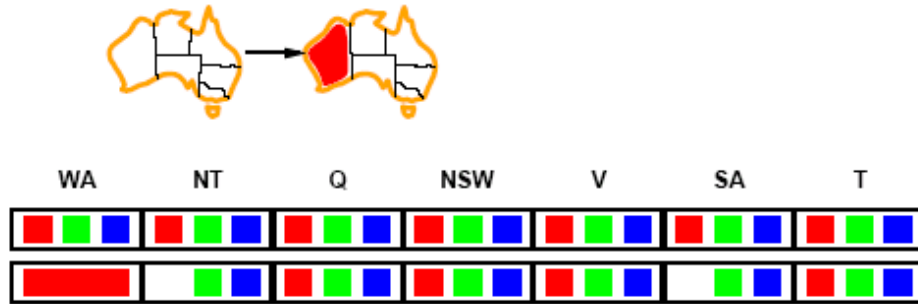
- Least constraining value heuristic
- Heuristic Rule: given a variable choose the least constraining value
 - leaves the maximum flexibility for subsequent variable assignments

Forward checking



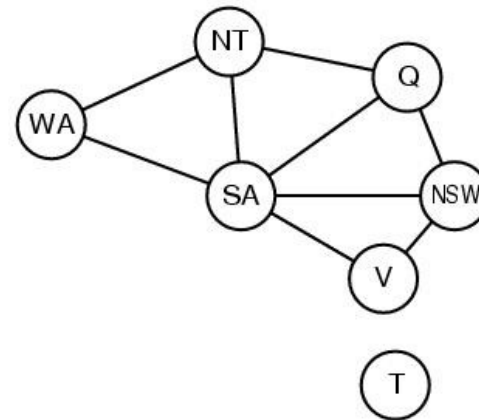
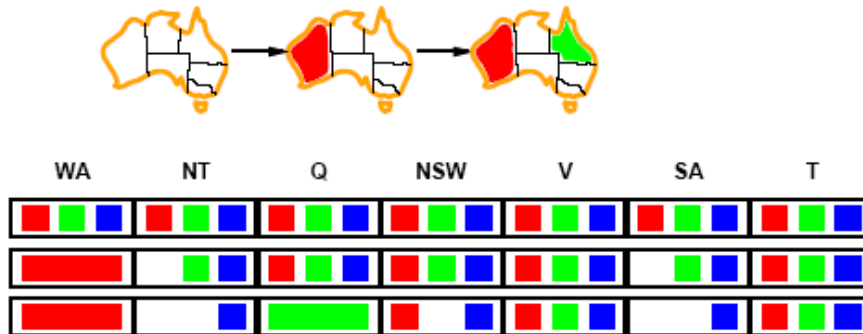
- Can we detect inevitable failure early?
 - *And avoid it later?*
- *Forward checking idea:* keep track of remaining legal values for unassigned variables.
- Terminate search when any variable has no legal values.

Forward checking



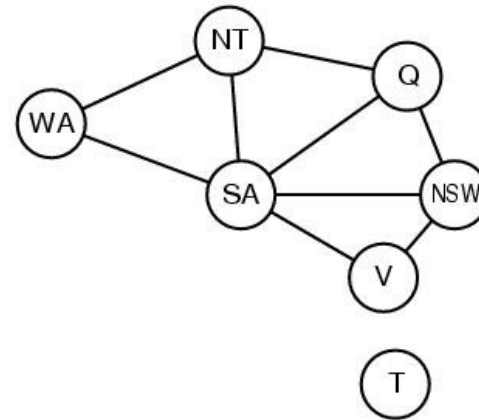
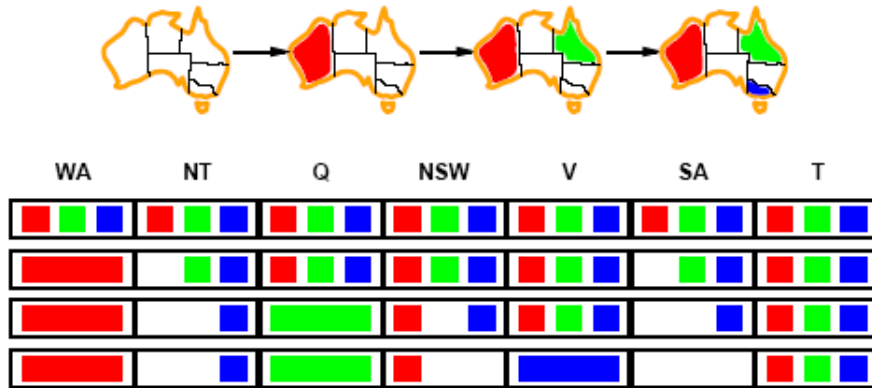
- Assign $\{WA=red\}$
- Effects on other variables connected by constraints to WA
 - *NT can no longer be red*
 - *SA can no longer be red*

Forward checking



- Assign $\{Q=green\}$
- Effects on other variables connected by constraints with WA
 - *NT can no longer be green*
 - *NSW can no longer be green*
 - *SA can no longer be green*
- *MRV heuristic* would automatically select NT or SA next

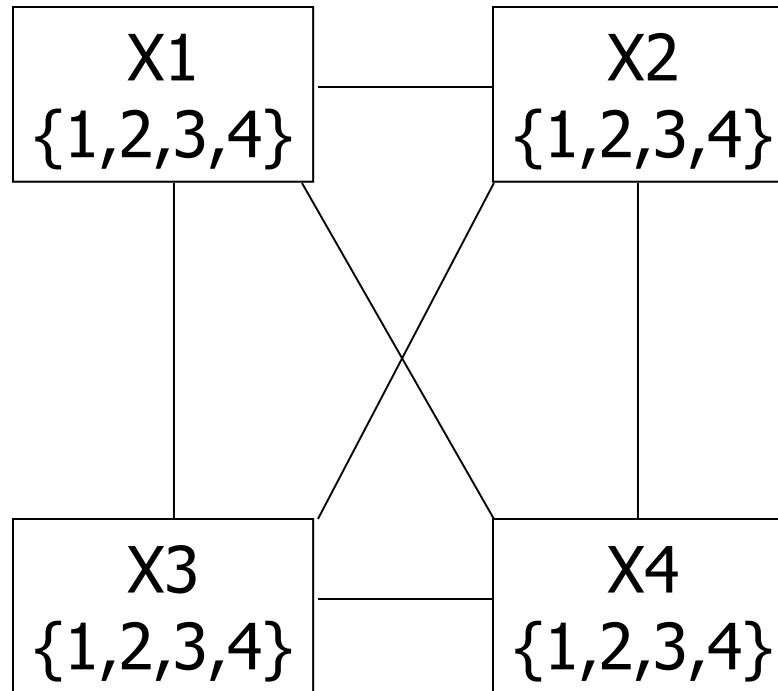
Forward checking



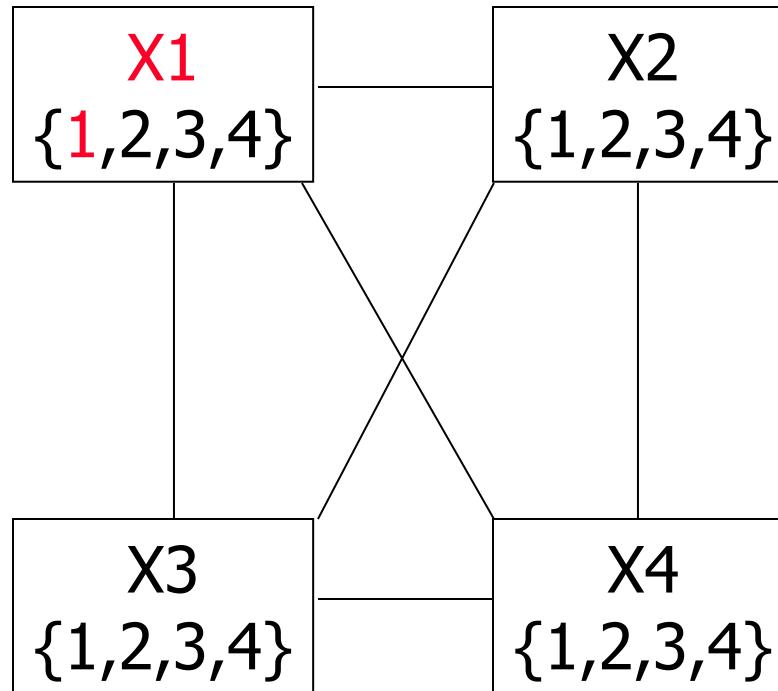
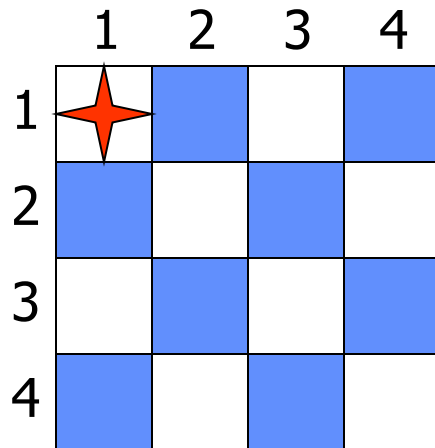
- If *V* is assigned *blue*
- Effects on other variables connected by constraints with *WA*
 - *NSW* can no longer be *blue*
 - *SA* is empty
- FC has detected that partial assignment is *inconsistent* with the constraints and backtracking can occur.

Example: 4-Queens Problem

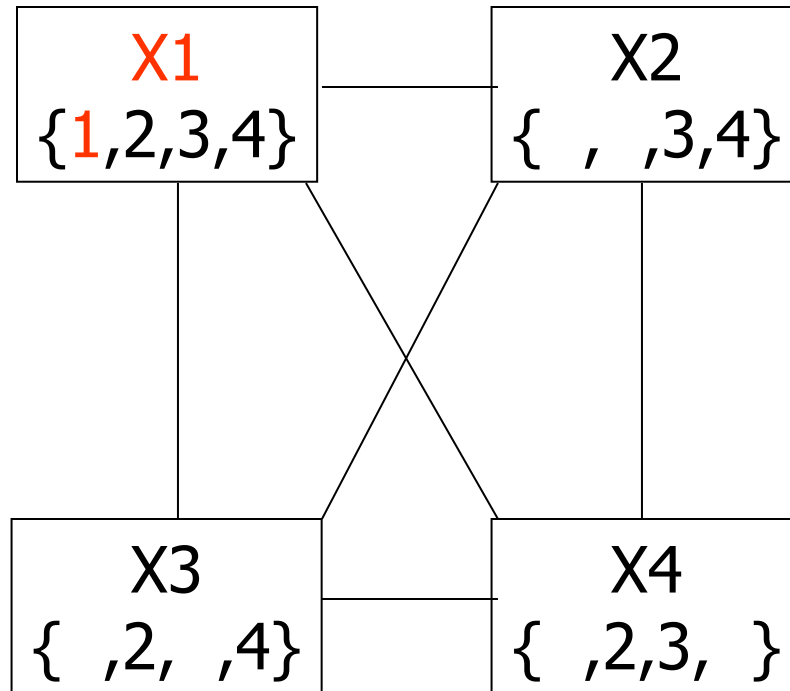

	1	2	3	4
1				
2				
3				
4				




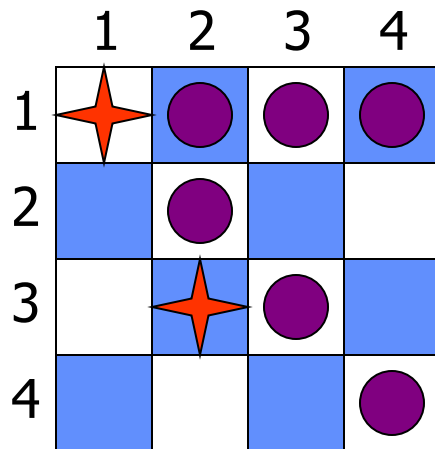
Example: 4-Queens Problem



Example: 4-Queens Problem

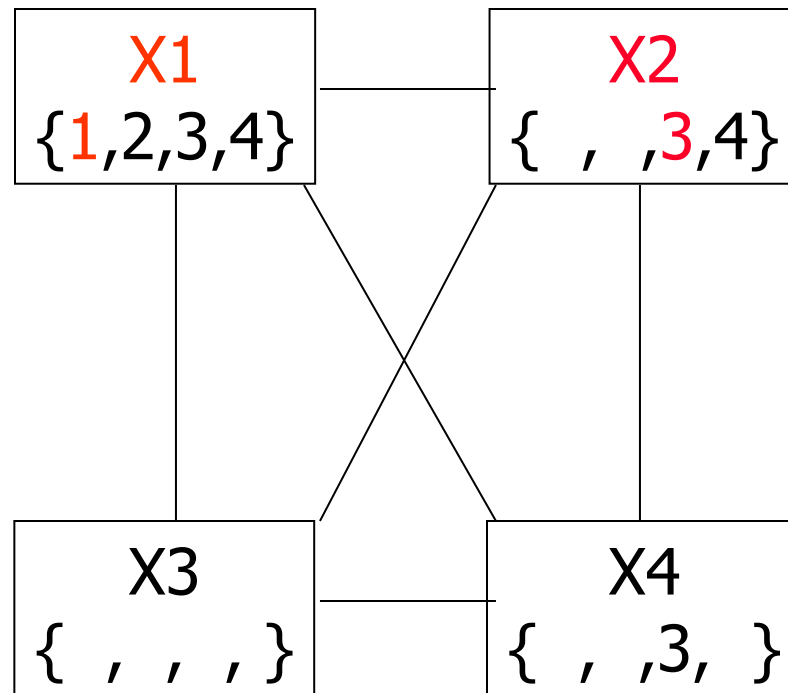


Example: 4-Queens Problem

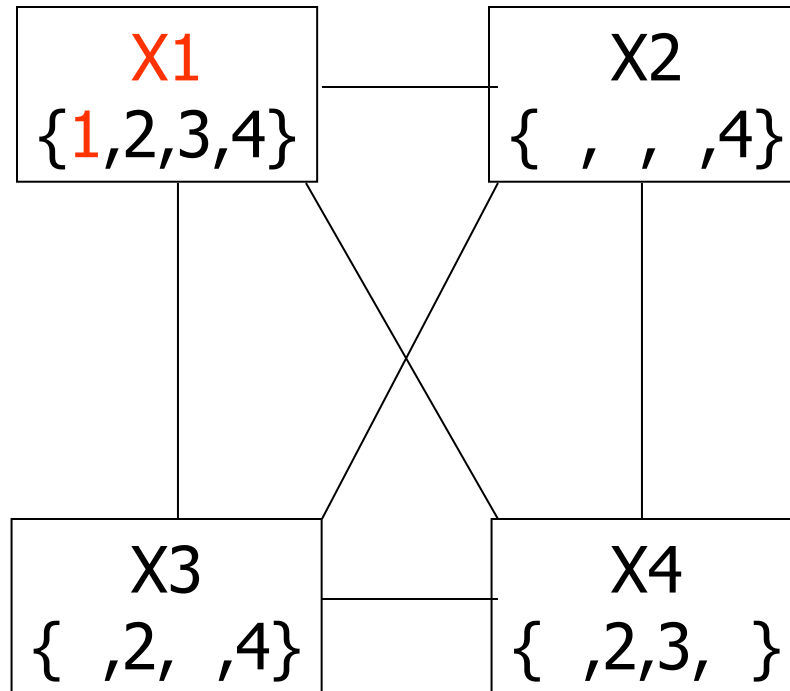
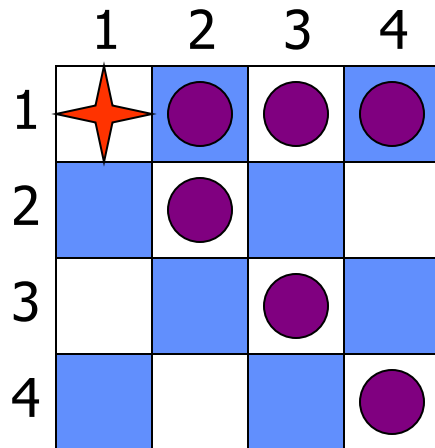


Example: 4-Queens Problem

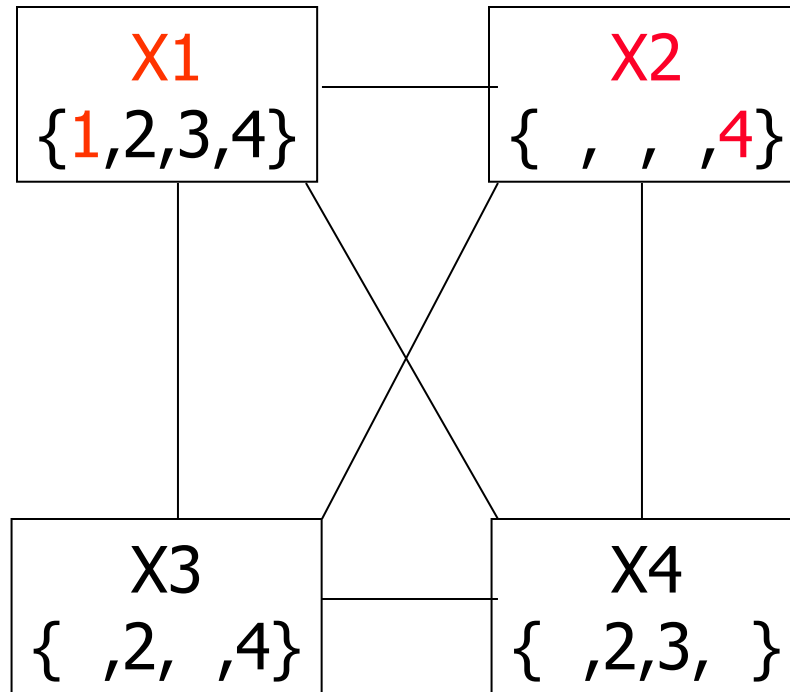
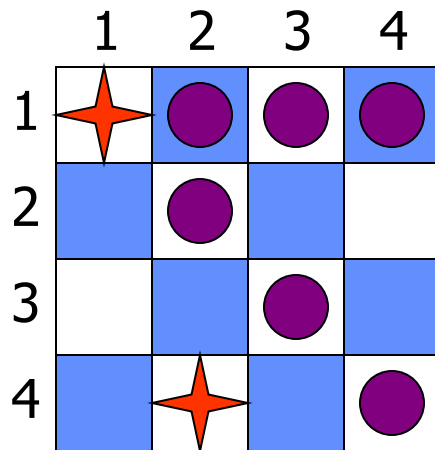
	1	2	3	4
1	★	●	●	●
2		●	●	
3		★	●	●
4			●	●



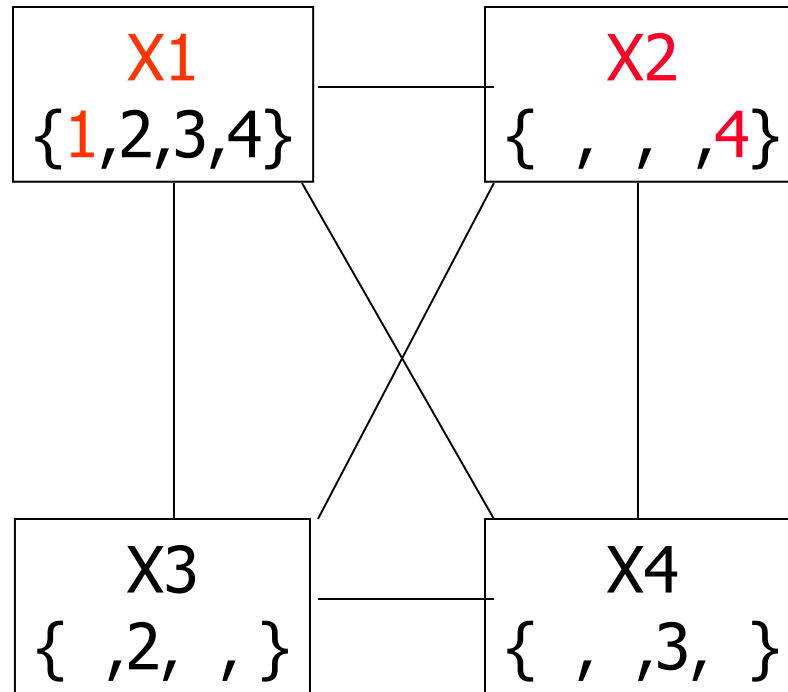
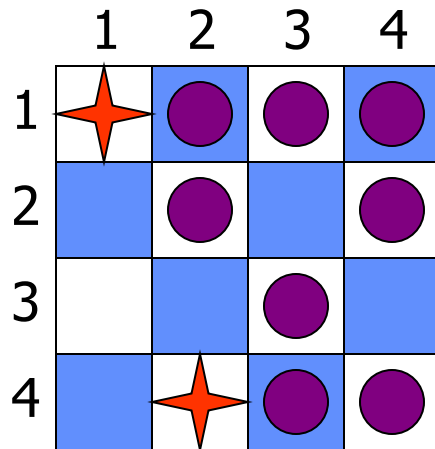
Example: 4-Queens Problem



Example: 4-Queens Problem




Example: 4-Queens Problem



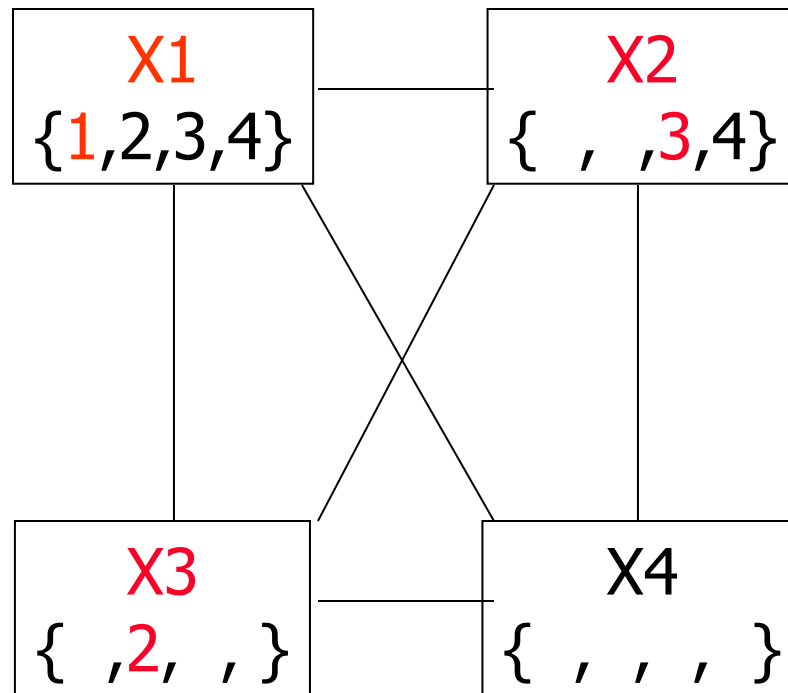
Example: 4-Queens Problem

	1	2	3	4
1	★	●	●	●
2		●	★	●
3			●	
4		★	●	●



Example: 4-Queens Problem

	1	2	3	4
1	★	●	●	●
2		●	★	●
3			●	●
4		★	●	●



Comparison of CSP algorithms on different problems

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA	(> 1,000K)	(> 1,000K)	2K	60	64
<i>n</i> -Queens	(> 40,000K)	13,500K	(> 40,000K)	817K	4K
Zebra	3,859K	1K	35K	0.5K	2K
Random 1	415K	3K	26K	2K	
Random 2	942K	27K	77K	15K	

Median number of consistency checks over 5 runs to solve problem

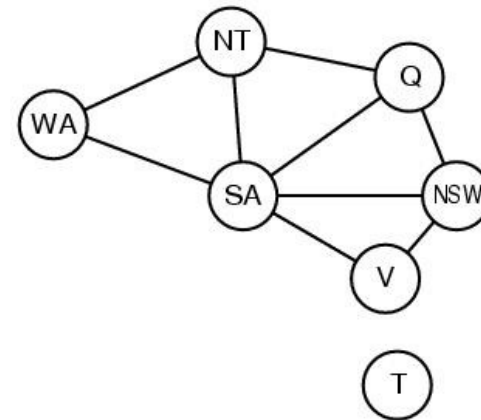
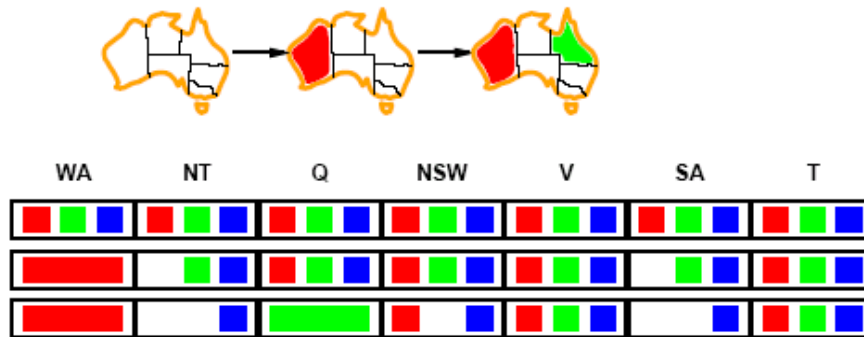
Parentheses -> no solution found

USA: 4 coloring

n-queens: $n = 2$ to 50

Zebra: see exercise 5.13

Constraint propagation

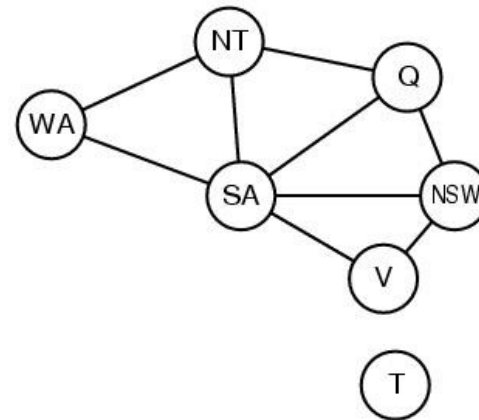
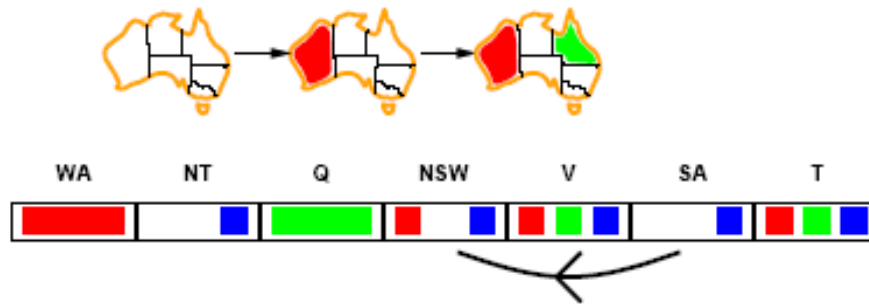


- Solving CSPs with combination of heuristics plus forward checking is more efficient than either approach alone
- FC checking does not detect all failures.
 - E.g., NT and SA cannot be blue

Constraint propagation

- Techniques like CP and FC are in effect eliminating parts of the search space
 - Somewhat complementary to search
- Constraint propagation goes further than FC by repeatedly enforcing constraints locally
 - Needs to be faster than actually searching to be effective
- Arc-consistency (AC) is a systematic procedure for constraint propagation

Arc consistency




- An Arc $X \rightarrow Y$ is consistent if
for every value x of X there is some value y consistent with x
(note that this is a directed property)

- Consider state of search after WA and Q are assigned:

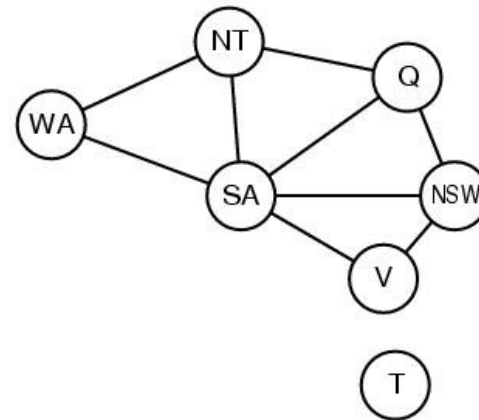
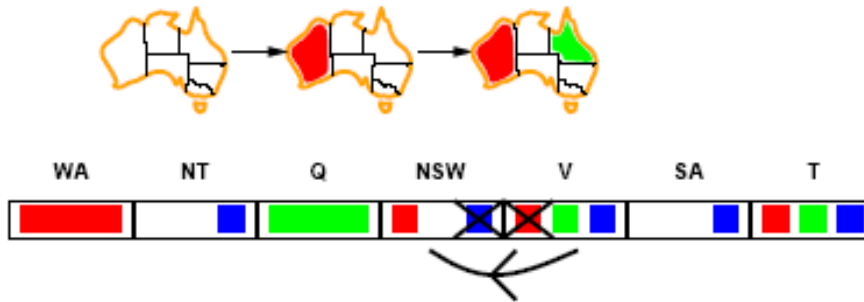
$SA \rightarrow NSW$ is consistent if
 $SA=blue$ and $NSW=red$

Arc consistency





- $X \rightarrow Y$ is consistent if
for every value x of X there is some value y consistent with x
- $NSW \rightarrow SA$ is consistent if
NSW=red and SA=blue
NSW=blue and SA=???

Arc consistency



- Can enforce arc-consistency:
Arc can be made consistent by removing *blue* from *NSW*
- Continue to propagate constraints....
 - Check $V \rightarrow NSW$
 - Not consistent for $V = \text{red}$
 - Remove red from V



- Continue to propagate constraints....
- $SA \rightarrow NT$ is not consistent
 - and cannot be made consistent
- Arc consistency detects failure earlier than FC

Arc consistency checking

- Can be run as a preprocessor or after each assignment
 - Or as preprocessing before search starts
- AC must be run repeatedly until no inconsistency remains
- Trade-off
 - Requires some overhead to do, but generally more effective than direct search
 - In effect it can eliminate large (inconsistent) parts of the state space more effectively than search can
- Need a systematic method for arc-checking
 - If X loses a value, neighbors of X need to be rechecked:
i.e. incoming arcs can become inconsistent again
(outgoing arcs will stay consistent).

Arc consistency algorithm (AC-3)

function AC-3(*csp*) **returns** false if inconsistency found, else true, may reduce *csp* domains

inputs: *csp*, a binary CSP with variables $\{X_1, X_2, \dots, X_n\}$

local variables: *queue*, a queue of arcs, initially all the arcs in *csp*

/ initial queue must contain both (X_i, X_j) and (X_j, X_i) */*

while *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$

if REMOVE-INCONSISTENT-VALUES(X_i, X_j) **then**

if size of $D_i = 0$ **then return** false

for each X_k **in** NEIGHBORS[X_i] – $\{X_j\}$ **do**

add (X_k, X_i) to *queue* if not already there

return true

function REMOVE-INCONSISTENT-VALUES(X_i, X_j) **returns** true iff we delete a value from the domain of X_i

removed \leftarrow false

for each x **in** DOMAIN[X_i] **do**

if no value y in DOMAIN[X_j] allows (x, y) to satisfy the constraints between X_i and X_j

then delete x from DOMAIN[X_i]; *removed* \leftarrow true

return *removed*

(from Mackworth, 1977)

Complexity of AC-3

- A binary CSP has at most n^2 arcs
- Each arc can be inserted in the queue d times (worst case)
 - (X, Y) : only d values of X to delete
- Consistency of an arc can be checked in $O(d^2)$ time
- Complexity is $O(n^2 d^3)$
- Although substantially more expensive than Forward Checking, Arc Consistency is usually worthwhile.

K-consistency

- Arc consistency does not detect all inconsistencies:
 - Partial assignment $\{WA=red, NSW=red\}$ is inconsistent.
- Stronger forms of propagation can be defined using the notion of k-consistency.
- A CSP is k-consistent if for any set of k-1 variables and for any consistent assignment to those variables, a consistent value can always be assigned to any kth variable.
 - E.g. 1-consistency = node-consistency
 - E.g. 2-consistency = arc-consistency
 - E.g. 3-consistency = path-consistency
- Strongly k-consistent:
 - k-consistent for all values $\{k, k-1, \dots, 2, 1\}$

Trade-offs

- Running stronger consistency checks...
 - Takes more time
 - But will reduce branching factor and detect more inconsistent partial assignments
 - No “free lunch”
 - In worst case n -consistency takes exponential time
- Generally helpful to enforce 2-Consistency (Arc Consistency)
- Sometimes helpful to enforce 3-Consistency
- Higher levels may take more time to enforce than they save.

Further improvements

- Checking special constraints
 - Checking Alldif(...) constraint
 - *E.g. {WA=red, NSW=red}*
 - Checking Atmost(...) constraint
 - *Bounds propagation for larger value domains*
- Intelligent backtracking
 - Standard form is chronological backtracking i.e. try different value for preceding variable.
 - More intelligent, backtrack to conflict set.
 - Set of variables that caused the failure or set of previously assigned variables that are connected to X by constraints.
 - Backjumping moves back to most recent element of the conflict set.
 - Forward checking can be used to determine conflict set.

Local search for CSPs

- Use complete-state representation
 - Initial state = all variables assigned values
 - Successor states = change 1 (or more) values
- For CSPs
 - allow states with unsatisfied constraints (unlike backtracking)
 - operators **reassign** variable values
 - hill-climbing with n-queens is an example
- Variable selection: randomly select any conflicted variable
- Value selection: *min-conflicts heuristic*
 - Select new value that results in a minimum number of conflicts with the other variables

Local search for CSP

function MIN-CONFLICTS(*csp*, *max_steps*) **return** solution or failure

inputs: *csp*, a constraint satisfaction problem

max_steps, the number of steps allowed before giving up

current \leftarrow an initial complete assignment for *csp*

for *i* = 1 to *max_steps* **do**

if *current* is a solution for *csp* then **return** *current*

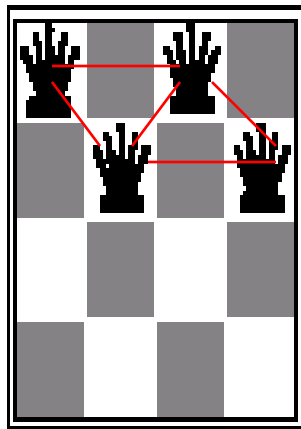
var \leftarrow a randomly chosen, conflicted variable from VARIABLES[*csp*]

value \leftarrow the value *v* for *var* that minimize CONFLICTS(*var*, *v*, *current*, *csp*)

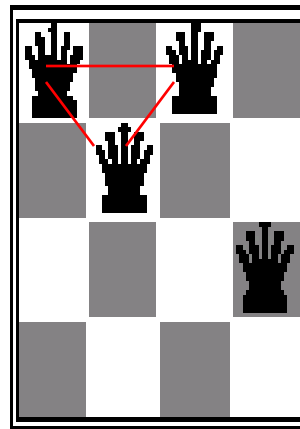
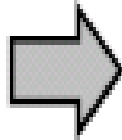
 set *var* = *value* in *current*

return *failure*

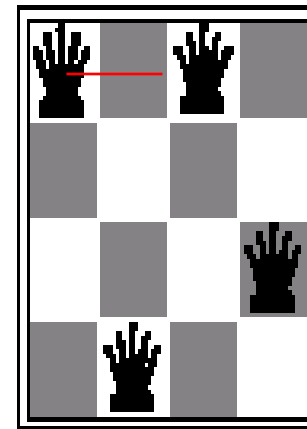
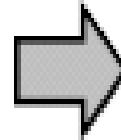
Min-conflicts example 1



$h=5$



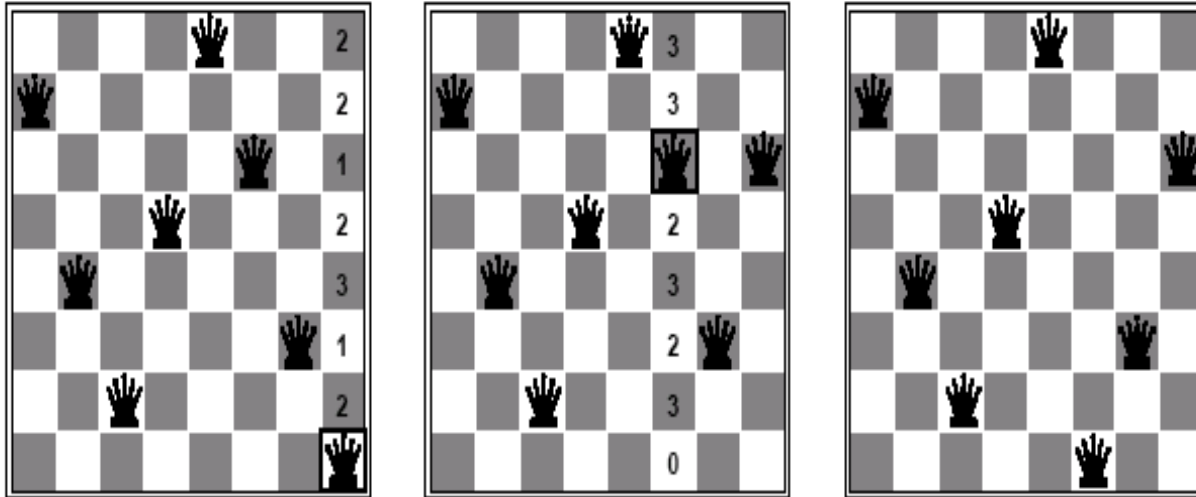
$h=3$



$h=1$

Use of min-conflicts heuristic in hill-climbing.

Min-conflicts example 2



- A two-step solution for an 8-queens problem using min-conflicts heuristic
- At each stage a queen is chosen for reassignment in its column
- The algorithm moves the queen to the min-conflict square breaking ties randomly.

Comparison of CSP algorithms on different problems

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA	(> 1,000K)	(> 1,000K)	2K	60	64
<i>n</i> -Queens	(> 40,000K)	13,500K	(> 40,000K)	817K	4K
Zebra	3,859K	1K	35K	0.5K	2K
Random 1	415K	3K	26K	2K	
Random 2	942K	27K	77K	15K	

Median number of consistency checks over 5 runs to solve problem

Parentheses -> no solution found

USA: 4 coloring

n-queens: $n = 2$ to 50

Zebra: see exercise 6.7 (3rd ed.); exercise 5.13 (2nd ed.)

Advantages of local search

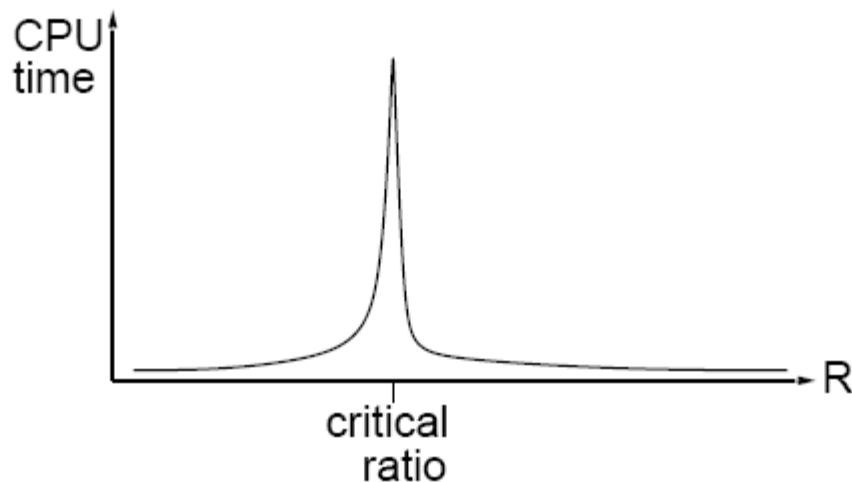
- Local search can be particularly useful in an online setting
 - Airline schedule example
 - E.g., mechanical problems require than 1 plane is taken out of service
 - Can locally search for another “close” solution in state-space
 - Much better (and faster) in practice than finding an entirely new schedule
- The runtime of min-conflicts is roughly independent of problem size.
 - Can solve the millions-queen problem in roughly 50 steps.
 - Why?
 - n-queens is easy for local search because of the relatively high density of solutions in state-space

Performance of min-conflicts

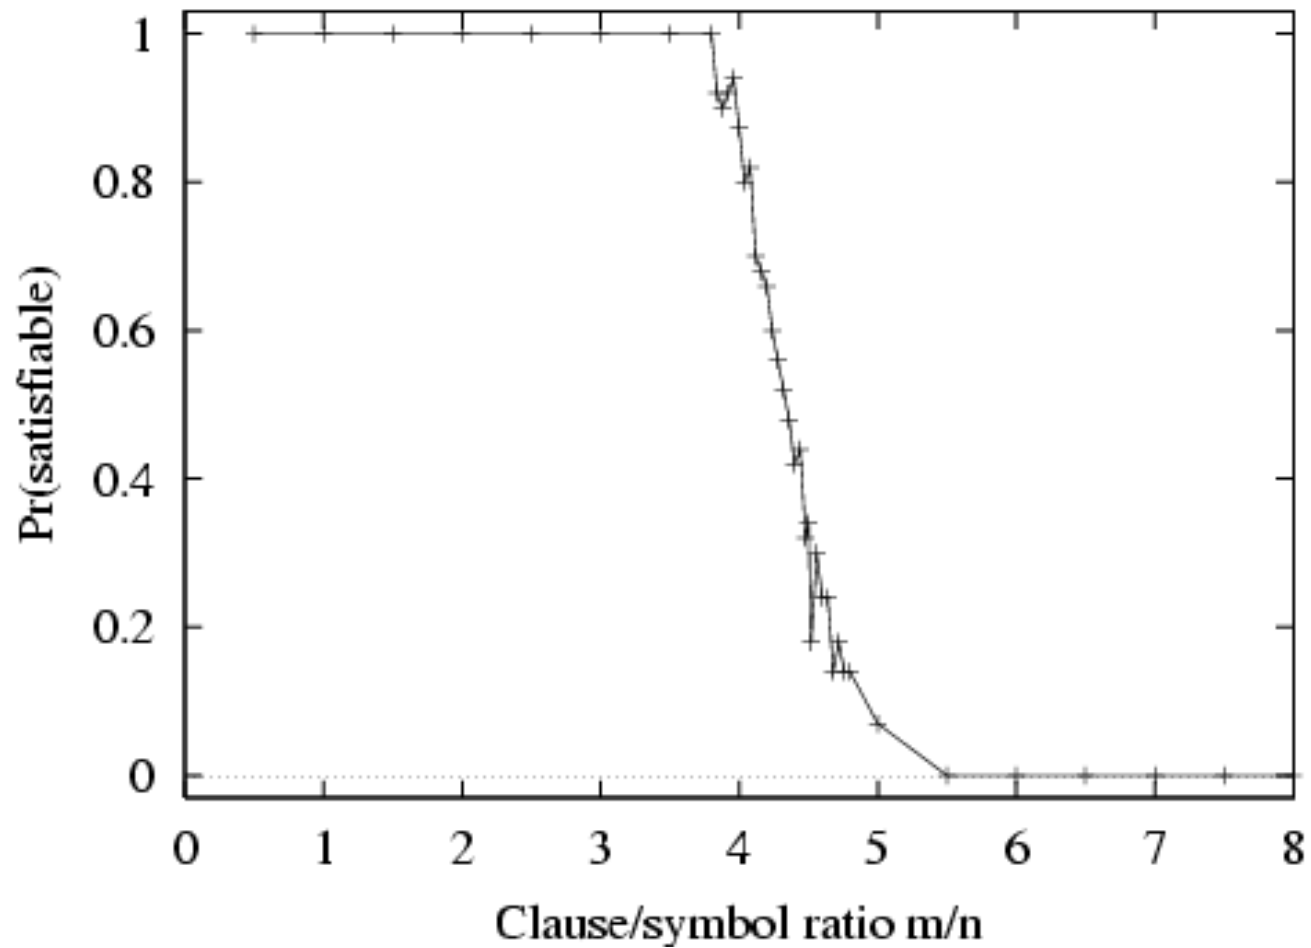
Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)

The same appears to be true for any randomly-generated CSP
except in a narrow range of the ratio


$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



Hard satisfiability problems

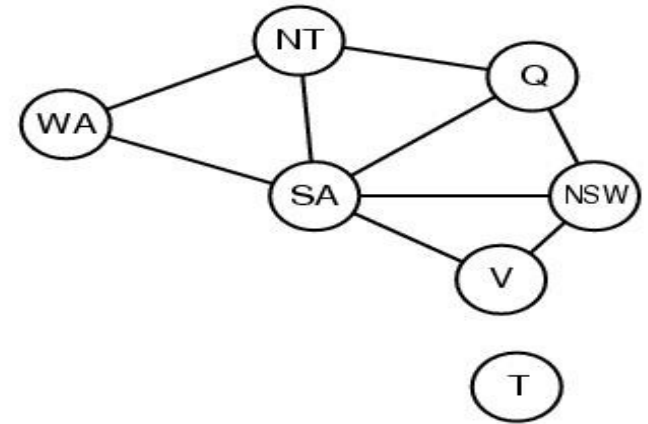


Hard satisfiability problems




- Median runtime for 100 **satisfiable** random 3-CNF sentences, $n = 50$

Graph structure and problem complexity



- Solving disconnected subproblems
 - Suppose each subproblem has c variables out of a total of n .
 - Worst case solution cost is $O(n/c d^c)$, i.e. linear in n
 - Instead of $O(d^n)$, exponential in n
- E.g. $n=80, c=20, d=2$
 - $2^{80} = 4$ billion years at 1 million nodes/sec.
 - $4 * 2^{20} = .4$ second at 1 million nodes/sec

Tree-structured CSPs



- Theorem:
 - if a constraint graph has no loops then the CSP can be solved in $O(nd^2)$ time
 - linear in the number of variables!
- Compare difference with general CSP, where worst case is $O(d^n)$

Summary

- CSPs
 - special kind of problem: states defined by values of a fixed set of variables, goal test defined by constraints on variable values
- Backtracking=depth-first search with one variable assigned per node
- Heuristics
 - Variable ordering and value selection heuristics help significantly
- Constraint propagation does additional work to constrain values and detect inconsistencies
 - Works effectively when combined with heuristics
- Iterative min-conflicts is often effective in practice.
- Graph structure of CSPs determines problem complexity
 - e.g., tree structured CSPs can be solved in linear time.